

Assigned Project

Dr. Min Chi

Special Thanks to: Shitian Shen
Department of Computer Science
North Carolina State University

Introduction

- Intelligent Tutoring System contains a set of actions
- Deep Thought (Dr. Barnes, 2015) can take two actions:
 - *Problem Solving* (PS)
 - *Work Example* (WE)

Problem Solving

1: $A \rightarrow (B \wedge C)$ 2: $A \vee D$ 3: $\neg D \wedge E$

Level: 1/6
Problem: 1/3

Problem Code: 1.0.1.0

C: B

Rules

MP ?
Modus Ponens

MT ?
Modus Tollens

DS ?
Disjunctive Syllogism

Add ?
Addition

Simp ?
Simplification

Conj ?
Conjunction

Hypothetical Syllogism

$p \rightarrow q$
 $q \rightarrow r$
 $\swarrow \searrow$
 $p \rightarrow r$

$q \rightarrow r$
 $p \rightarrow q$
 $\swarrow \searrow$
 $p \rightarrow r$

Commutative

Assocative

Dist ?
Distributive

Abs ?
Absorption

Exp ?
Exportation

Taut ?
Tautology

	Expression	Antecedent Lines	Rule Used
1	$A \rightarrow (B \wedge C)$		Given
2	$A \vee D$		Given
3	$\neg D \wedge E$		Given

For	Type	For	Type
$A \wedge B$	$A * B$	$A \rightarrow B$	$A > B$
$A \vee B$	$A + B$	$A \leftrightarrow B$	$A = B$
$\neg A$	$\neg A$		

Message Box

No blocks selected. Rule requires two justified premises

Representation

☒ Symbolic ☐ English

Delete Node

Change to Indirect Proof

Restart Current Problem

Skip Current Problem

Deep Thought

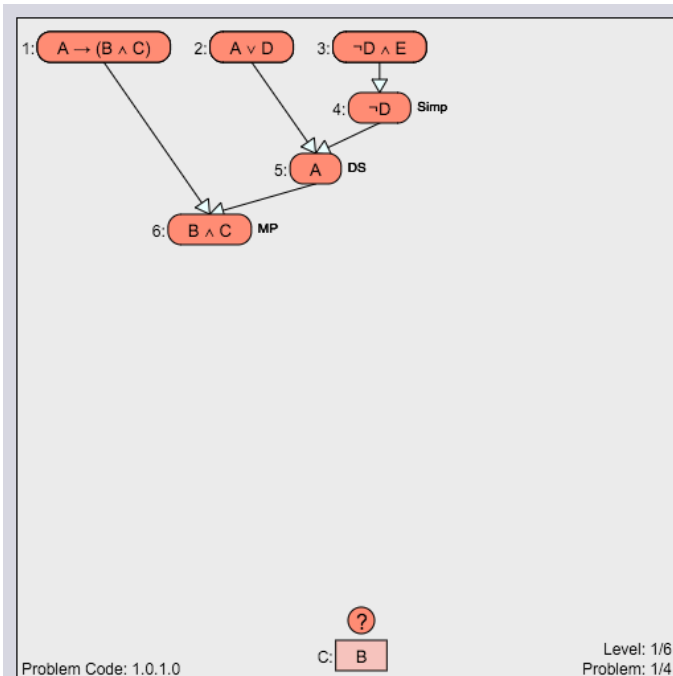
A Logic Proof Tutor
Version 2.0
August 15, 2014
North Carolina State University

Instructions

Window Information

Contact/Version Information

Work Example



Extract B from 6 using Simp

Message Box

EXAMPLE

Click the arrow next to the Hint Box to step through the example.

Representation

☒ Symbolic ☐ English

	Expression	Antecedent Lines	Rule Used
1	$A \rightarrow (B \wedge C)$		Given
2	$A \vee D$		Given
3	$\neg D \wedge E$		Given
4	$\neg D$	3	Simplification
5	A	2, 4	Disjunctive Syllogism
6	$B \wedge C$	1, 5	Modus Ponens

C B

For

$A \wedge B$

$A \vee B$

$\neg A$

Type

$A \wedge B$

$A \vee B$

$\neg A$

For

$A \rightarrow B$

$A \leftrightarrow B$

$A = B$

Type

$A \wedge B$

$A \vee B$

$A = B$

Deep Thought
A Logic Proof Tutor

Version 6

January 19, 2016

North Carolina State University

Instructions

Window Information

Contact/Version Information

Question

When to assign PS or WE to students ?

Pedagogical strategy is defined as policies to decide what the system action to take next in the face of alternatives.

Induce Pedagogical Strategy

- Inducing pedagogical strategy is challenging
 - Hard code
 - Data driven


Reinforcement Learning vs. Inducing Pedagogical strategy

What is the best action for the **agent** (*tutor*)
to take in any **state** (*learning context*)
in order to maximize **reward** (*student learning*)

Reinforcement Learning:

- Model-based vs Model-free Reinforcement Learning
 - Model-based
 - Generating data is expensive (ITS)
 - Learn from the model instead of data sets
 - Model-free
 - Collecting data is trivial (playing chess)
 - Learn from data sets directly

Agent Environment Interaction



Markov Decision Process: Definition

- A Mathematical framework for representing a reinforcement learning task
- A tuple $\langle S, A, T, R, \pi \rangle$

State Set	
Action Set	
Transition Probability	
Rewards	
Policy	

Value Iteration: Algorithm

1. $V_0(s) = 0, \text{ for } s \in S$

Initialization

2. For k

$\Delta \leftarrow 0$

For each $s \in S$

$v \leftarrow V_{k-1}(s)$

$V_k(s) \leftarrow \max_a \sum_{s'} T_{ss'}^a [R_{ss'}^a + \gamma V_{k-1}(s')]$

$\Delta \leftarrow \max(\Delta, |v - V_k(s)|)$

Until $\Delta \leftarrow \theta$ (a small positive number)

Maximizing
Value Function

3. $\pi(s) = \underset{a}{\operatorname{argmax}} \sum_{s'} T_{ss'}^a [R_{ss'}^a + \gamma V^\pi(s')]$

Policy
Generation

Value Iteration: Example

- Transfer Data into trajectories
 - State set : $\{S_1, S_2\}$
 - Action set: $\{PS, WE\}$

$$\begin{array}{cccccccccccc}
 S_1 & \xrightarrow{PS, 0} & S_2 & \xrightarrow{PS, 0} & S_2 & \xrightarrow{WE, 50} & S_1 & \xrightarrow{PS, 0} & \dots & \xrightarrow{WE, 0} & S_2 & \xrightarrow{PS, 100} & S_1 & \xrightarrow{WE, 0} \\
 S_2 & \xrightarrow{PS, 0} & S_2 & \xrightarrow{WE, 50} & S_1 & \xrightarrow{PS, 0} & \dots & \xrightarrow{WE, 0} & S_2 & \xrightarrow{PS, 0} & S_1 & \xrightarrow{WE, -50} & S_1 & \xrightarrow{WE, 0} \\
 S_2 & \xrightarrow{PS, 100} & S_1 & \xrightarrow{WE, 0} & S_2 & \xrightarrow{PS, 100} & \dots & \xrightarrow{WE, 0} & S_1 & \xrightarrow{PS, 0} & S_2 & \xrightarrow{WE, 0} & S_2 & \xrightarrow{WE, 0} T
 \end{array}$$

Value Iteration: Example

- Transition Probability

$$P(S_1|S_2, PS) = \frac{\#(S_2 \xrightarrow{PS} S_1)}{\#(S_2 \xrightarrow{PS} S_1) + \#(S_2 \xrightarrow{PS} S_2)} = \frac{1}{4}$$

- Expected Rewards

$$R(S_1|S_2, PS) = \frac{\sum r(S_2 \xrightarrow{PS} S_1)}{\#(S_2 \xrightarrow{PS} S_1)} = 20$$

Value Iteration: Example

- Transition probability $T_{ss'}^a$

PS

1/4	3/4
1/2	1/2

WE

1/2	1/2
2/3	1/3

- Reward function $R_{ss'}^a$

PS

10	40
20	30

WE

20	30
45	5

Value Iteration: Example

K			
0	0		0
1	32.50	PS	31.67 WE
2	61.18	PS	60.67 WE
3	87.22	PS	86.58 WE
4	110.56	PS	109.97 WE
121	320.90	PS	320.30 WE
122	320.90	PS	320.30 WE

$$V_1(S_1) = \max \left\{ \begin{array}{l} \frac{1}{4}(10 + 0.9 * 0) + \frac{3}{4}(40 + 0.9 * 0) = 32.50 \quad PS \\ 1 \end{array} \right.$$

$$V_1(S_2) = \max \left\{ \begin{array}{l} \frac{1}{2}(20 + 0.9 * 0) + \frac{1}{2}(30 + 0.9 * 0) = 25 \quad PS \\ \frac{2}{5}(45 + 0.9 * 0) + \frac{1}{5}(5 + 0.9 * 0) = 31.67 \quad WE \end{array} \right.$$

$$V_2(S_1) = \max \left\{ \begin{array}{l} \frac{1}{4}(10 + 0.9 * 32.5) + \frac{3}{4}(40 + 0.9 * 31.67) = 61.18 \quad PS \\ 1 \end{array} \right.$$

$$V_2(S_2) = \max \left\{ \begin{array}{l} \frac{1}{2}(20 + 0.9 * 32.5) + \frac{1}{2}(30 + 0.9 * 31.67) = 53.87 \quad PS \\ \frac{2}{3}(45 + 0.9 * 32.5) + \frac{1}{3}(5 + 0.9 * 31.67) = 60.67 \quad WE \end{array} \right.$$

Optimal policy π^* :

$S_1 \rightarrow PS$
 $S_2 \rightarrow WE$

Policy Evaluation

- Expected Cumulative Reward (Tetreault, 2006)

$$ECR = \sum_{i=1}^m \frac{N_i}{N_1 + N_2 + \dots + N_m} \times V^{\pi}(S_i)$$

Where S_i is the starting state, N_i is the times that S_i exists as starting state

- The higher ECR of the policy means the better policy

The Challenge is:

What is the best action for the **agent** (*tutor*)
to take in any **state** (*learning context*)
in order to maximize **reward** (*student learning*)

Challenge: State Representation

How to design states representing environment ?

State Representation: Feature Selection for RL

- Three types of feature selection methods
 - Filtered approach
 - Feature Selection process is independent to model construction
 - Evaluating the independence between reward with feature (Hirotsuka, Masashi 2010)
 - Wrapper approach
 - Feature subsets are evaluated by predefined score function
 - Monte Carlo tree search algorithm (Gaudel 2010)
 - Embedded approach
 - Feature selection and model construction are executed simultaneously
 - Least Square Temporal Difference with lasso regularized item (Kolter 2009)

Previous research:

Correlation-based Methods:

High vs Low

- When selecting features, should we select the feature that is most correlated (High) or uncorrelated (Low) to current optimal feature set ?
- In Supervised Learning, features with high correlation with labels are selected (C. Lee, 2010; L Yu & H Liu, 2003)
- In RL, the answer is not straightforward

Research Question: Low vs. High

- Choosing most correlated features (High)
 - Most likely to be related to decision making
 - May not make more contribute than current optimal feature set
- Choosing most uncorrelated features (Low)
 - Raise the diversity of feature set
 - Take the risk of involving irrelevant or noisy features

Correlation Metrics

Given labeled data, we can compute some simple score $S(i)$ that measures how informative each feature X is about the class labels Y .

- Chi-square (CHI) (Zibran, 2007)

$$\chi^2 = \sum_i \frac{(X_i - Y_i)^2}{Y_i}$$

- Information Gain (IG) (C. Lee, 2010)

$$IG(X, Y) = H(Y) - H(X|Y)$$

Correlation Metrics

- Information Gain Ratio (IGR) (J. T. Kent, 1983)

$$IGR(X, Y) = \frac{H(X) - H(X|Y)}{H(Y)}$$

- Symmetric Uncertainty (SU) (L. Yu, H. Liu, 2003)

$$SU(X, Y) = \frac{H(X) - H(X|Y)}{H(X) + H(Y)}$$

- Weighted Information Gain (WIG) (We proposed)

$$WIG(X, Y) = \frac{H(X) - H(X|Y)}{(H(X) + H(Y))H(Y)}$$

Correlation-based Feature Selection Methods

- Feature Selection for model-based RL
- Apply correlation between current optimal feature set and potential feature as the feature selection criteria
- Forward feature selection strategy

10 Correlation-based Methods

- Explore both high and low correlation
- Obtain 10 correlation-based feature selection methods (5 correlation metrics \times 2 correlation types)

	High	Low
CHI	CHI-High	CHI-Low
IG	IG-High	IG-Low
IGR	IGR-High	IGR-Low
SU	SU-High	SU-Low
WIG	WIG-High	WIG-Low

Other Implemented Methods

- Ensemble Methods
 - 10 correlation-based methods
 - 4 RL based methods
- RLPreviousFS
 - 4 RL based methods
 - 2 PCA based methods
 - 4 PCA & RL based methods

Intelligent Tutoring System

- Deep Thought (Dr. Barnes, 2015)
 - A rule-based tutoring system for teaching logic proof problems
 - Student solves 1-3 problems per level (Total 6 levels)
 - Level score ($LevelScore_i, i \in [1,6]$) is given for each student based on his/her performance on the last problem in the level i

Deep Thought : Reward Function

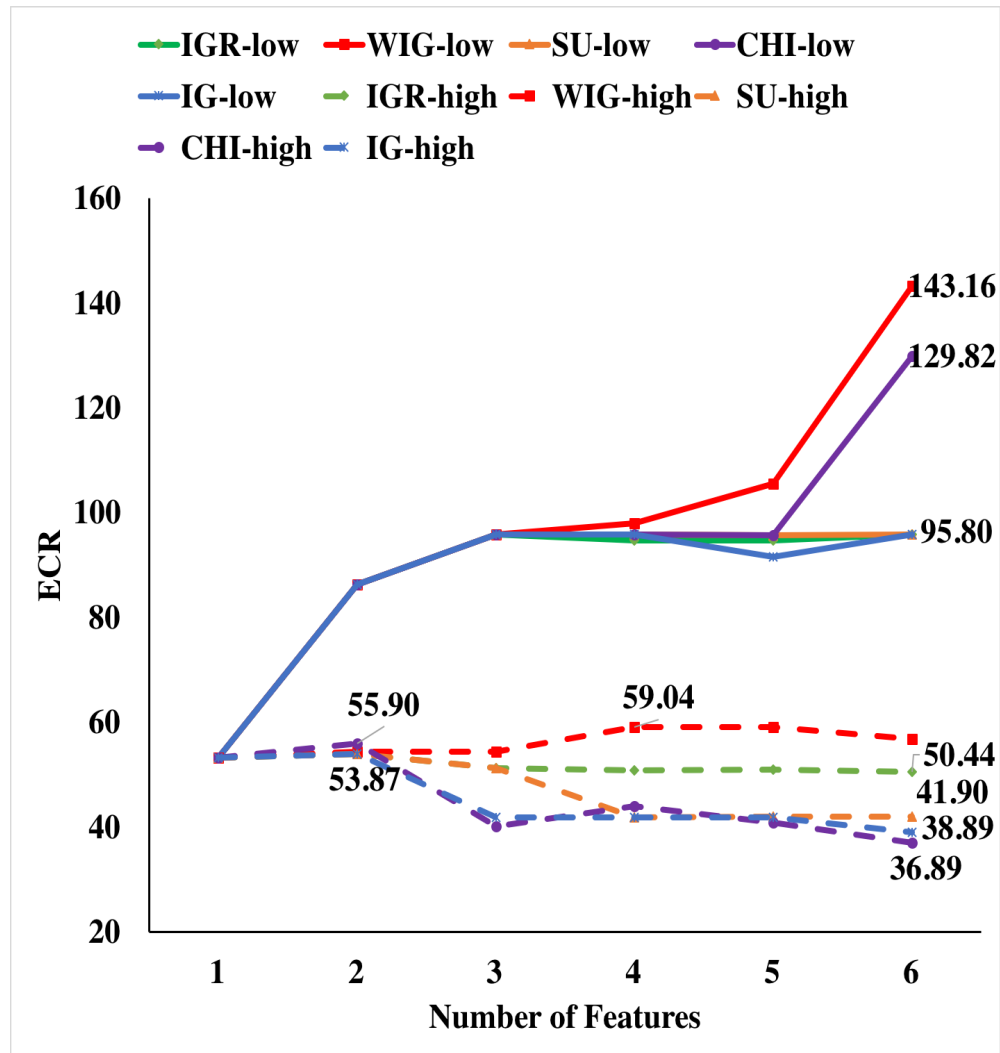
- Immediate Reward
 - $R_1 = LevelScore_1$
 - $R_i = LevelScore_i - LevelScore_{i-1}, i \in [2,6]$
- Delayed Reward

$$R_{delay} = LevelScore_6 - LevelScore_1$$

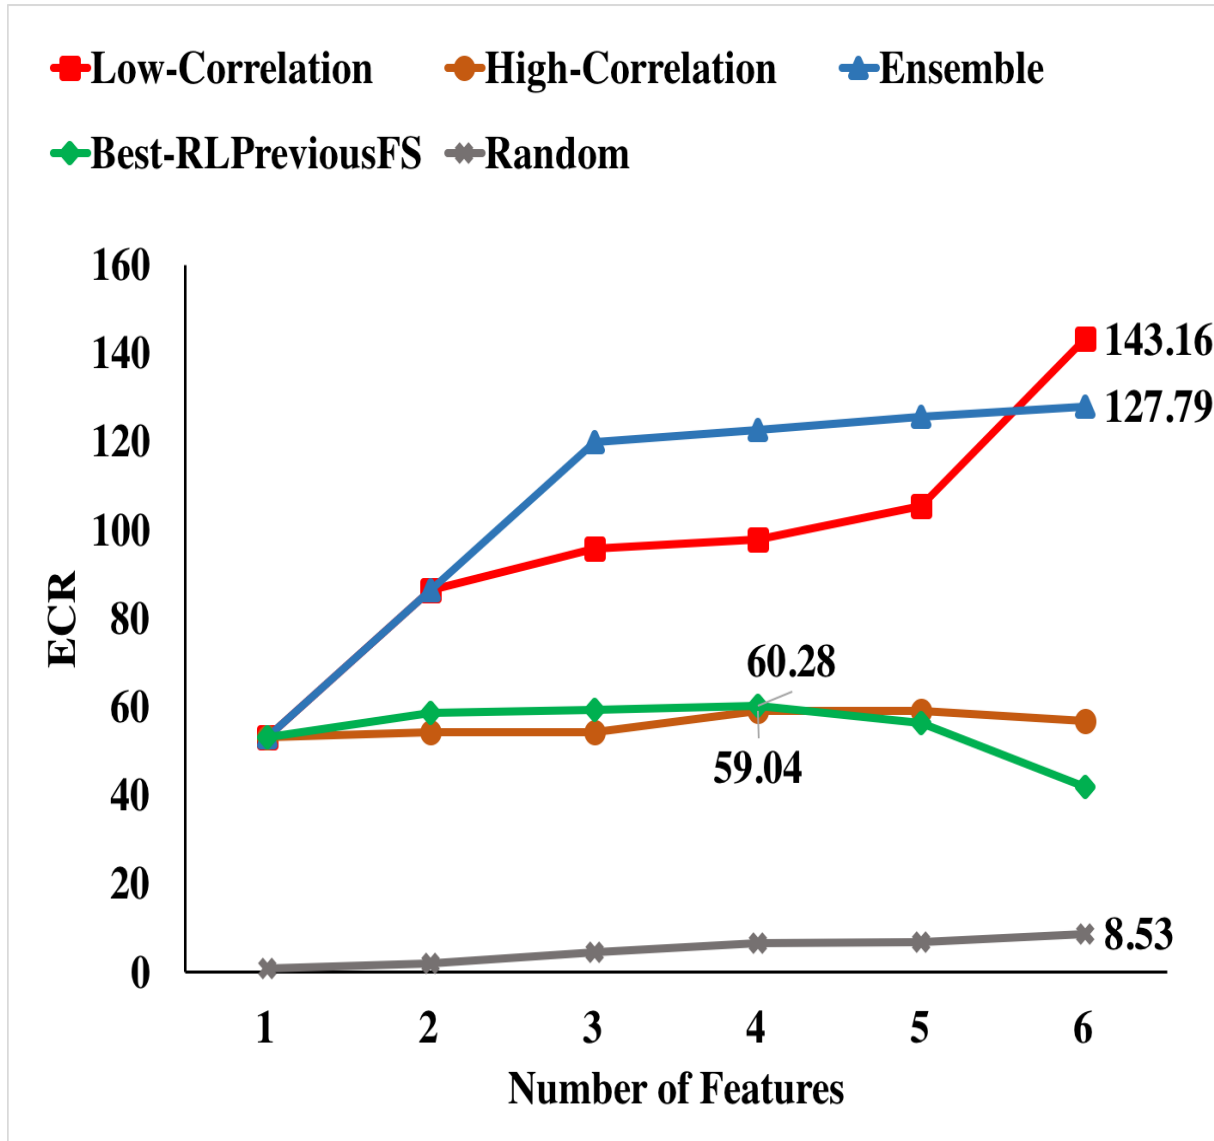
Deep Thought Data Sets

- Total 303 students in Fall 2014 and Spring 2015
- Average time spend in tutor is 416.60 minutes
- Total 135 features
- Action set
 - should it ask student to solve the next problem (PS)
 - should it provide an example to show the student how to solve the next problem (WE)

Result: High vs Low correlation



Results: Overall Evaluation



Induced Pedagogical Strategy

		Last three features $f_{I4}:f_{I5}:f_{I6}$							
		0:0:0	0:0:1	0:1:0	0:1:1	1:0:0	1:0:1	1:1:0	1:1:1
First three features $f_{I1}:f_{I2}:f_{I3}:f_{I4}$	0:0:0:0								
	0:0:0:1								
	0:0:1:0								
	0:0:1:1								
	0:1:0:0								
	0:1:0:1								
	0:1:1:0								
	0:1:1:1								
	1:0:0:0								
	1:0:0:1								
	1:0:1:0								
	1:0:1:1								
	1:1:0:0								
	1:1:0:1								
	1:1:1:0								
	1:1:1:1								

64 rules associated with WE (White)

21 rules associated with PS (Black)

43 no rules (Gray)

Induced Pedagogical Strategy

First three features $f_{I1}:f_{I2}:f_{I3}:f_{I4}$

Last three features $f_{I4}:f_{I5}:f_{I6}$

	0:0:0	0:0:1	0:1:0	0:1:1	1:0:0	1:0:1	1:1:0	1:1:1
0:0:0:0								
0:0:0:1								
0:0:1:0								
0:0:1:1								
0:1:0:0								
0:1:0:1								
0:1:1:0								
0:1:1:1								
1:0:0:0								
1:0:0:1								
1:0:1:0								
1:0:1:1								
1:1:0:0								
1:1:0:1								
1:1:1:0								
1:1:1:1								

The best Policy

64 rules associated with WE (White)

21 rules associated with PS (Black)

43 no rules (Gray)

First three features $f_{D1}:f_{D2}:f_{D3}$

Last three features $f_{D4}:f_{D5}:f_{D6}$

	0:0:0	0:0:1	0:1:0	0:1:1	1:0:0	1:0:1	1:1:0	1:1:1
0:0:0								
0:0:1								
0:1:0								
0:1:1								
0:2:0								
0:2:1								
1:0:0								
1:0:1								
1:1:0								
1:1:1								
1:2:0								
1:2:1								

Another Policy

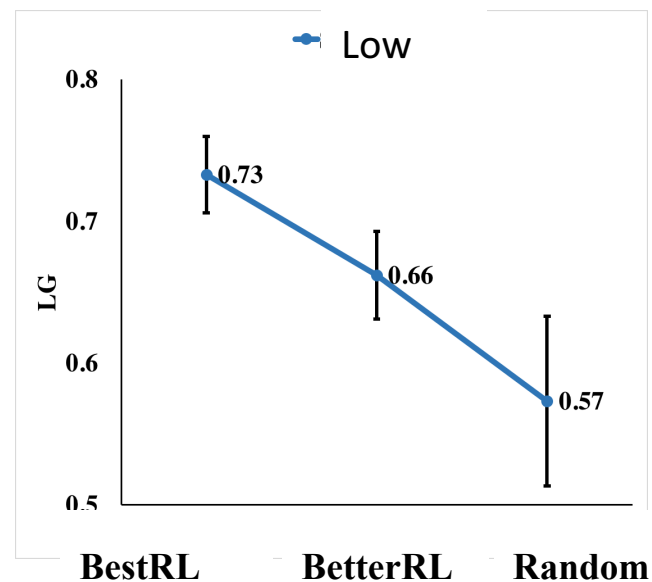
18 rules associated with WE

48 rules associated with PS

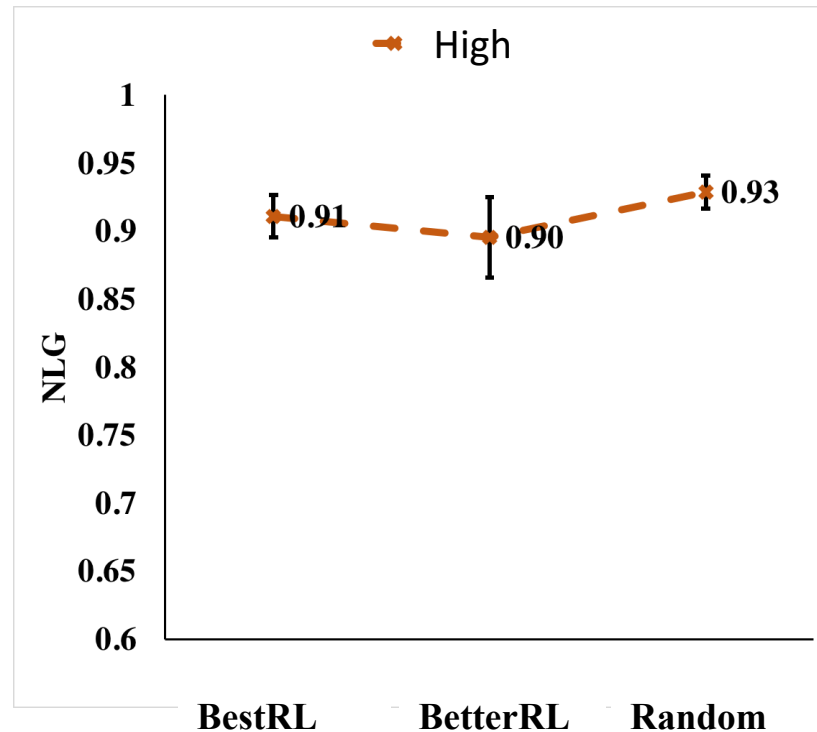
30 no rules

Learning Performance Result

- Significant difference among three Low groups
 - $F(2,46) = 3.99, p = 0.025$
- BestRLPolicy-Low group significant outperforms BetterRLPolicy-Low
 - $t(27) = 2.69, p = 0.012$
- BestRLPolicy-Low group marginally outperforms BetterRLPolicy-Low
 - $t(35) = 1.67, p = 0.098$




Learning Performance Result



- No significant difference between three High groups

Learning Performance Result



Your Task: State representation

- Discretization the features
- Feature Extraction and/or Feature selection (**explore new methods**)
- No more than 8 features (new or selected features)
- Evaluation: ECR
- Rank all the project: $[80-100] * 0.1$ points.
- Presentation: 5 points
- Submit your code and we will run it.

Publications

- Shitian Shen, M Chi, “*Reinforcement Learning: the Sooner the Better or the Later the Better ?*”, The 24th ACM User Modeling, Adaptation and Personalization (ACM UMAP), 2016 (Full paper)
- Shitian Shen, M Chi, “*Aim Low: Correlation-based Feature Selection for Model-based Reinforcement Learning*”, 9th International Conference on Educational Data Mining (EDM), 2016 (Short paper)
- Shitian Shen, M Chi, “*An Analysis of Feature Selection and Reward Function for Model-Based Reinforcement Learning*”, 13th International Conference on Intelligent Tutoring System (ITS), 2016 (Poster)